

Home Assignment: Sales Forecasting using Polynomial Regression

Learning Objective

The objective of this assignment is to apply **Polynomial Regression** to model a **non-linear relationship** between sales drivers and sales output.

Students will compare **Linear vs Polynomial Regression**, evaluate model performance, and justify model selection based on metrics and visualization.

Expected Completion Time

- **Best Case:** 60 minutes
 - **Average Case:** 90 minutes
-

Problem Statement

You are provided with a **Sales dataset** that captures the relationship between **Advertising Spend** and **Sales Revenue**.

Initial analysis shows that sales growth does not increase linearly with advertising spend.

Your task is to:

- Build a **Polynomial Regression model**
 - Compare it against a **Linear Regression baseline**
 - Decide which model better represents the data
-

Dataset Description

Dataset file: `sales_data.csv`

Column Name	Description
Advertising_Spend	Amount spent on advertising
Sales	Revenue generated

Assignment Requirements

Step 1: Data Exploration

- Load the dataset using Pandas
 - Display summary statistics
 - Plot a **scatter plot** between Advertising Spend and Sales
 - Comment on whether the relationship appears linear or non-linear
-

Step 2: Baseline Model – Simple Linear Regression

- Build a **Linear Regression** model using:
 - $X = \text{Advertising_Spend}$
 - $y = \text{Sales}$
 - Train and predict on the same dataset
 - Evaluate the model using:
 - Mean Squared Error (MSE)
 - R^2 Score
 - Plot the regression line over the scatter plot
-

Step 3: Polynomial Feature Transformation

- Use `PolynomialFeatures` to transform the independent variable
 - Train **Polynomial Regression models** with:
 - Degree 2
 - Degree 3
 - Clearly mention which degree you are using for final evaluation
-

Step 4: Polynomial Regression Model Training

- Train a Linear Regression model on the transformed polynomial features
 - Predict sales values using the polynomial model
-

Step 5: Model Evaluation and Comparison

Evaluate **both Linear and Polynomial models** using:

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- R^2 Score

Create a small comparison table showing:

- Model Type
 - MSE
 - R² Score
-

Step 6: Visualization

- Plot:
 - Actual sales data (scatter)
 - Linear regression line
 - Polynomial regression curve (smooth curve)
 - Ensure:
 - Clear labels
 - Legend
 - Proper title
-

Step 7: User Input Prediction

- Ask the user to enter a new **Advertising Spend** value
 - Predict sales using:
 - Linear Regression model
 - Polynomial Regression model
 - Print both predictions and briefly comment on the difference
-

Step 8: Model Interpretation (Important)

Answer the following questions in comments or markdown:

1. Why does Polynomial Regression perform better or worse than Linear Regression?
 2. What risks are associated with choosing a higher polynomial degree?
 3. In a real business scenario, which model would you choose and why?
-

Hints for Learners

- Polynomial Regression is still **Linear Regression under the hood**
- Sort X values before plotting polynomial curves for smooth visuals
- Use:
 - `np.argsort(X.flatten())`
- Higher R² alone does **not** guarantee a better model
- Watch out for **overfitting** when increasing degree